



Receiver cost and actual model usage

RTX 3060 · FP32 · batch size 1 | 4,000 distinct blocks per SNR, each timed three times

Mean end-to-end latency



Mean ByT5 invocations



A separate timing cohort, not the main BLER experiment. Repeated timings are not independent transmissions; early stopping reduces calls.